The Journal of Open Source Software

GEEODE: A Google Earth Engine Implementation of Optimization by Differential Evolution

Devin Routh [1] **and Claudia Roeoesli** [1]

1 Remote Sensing Laboratories, Department of Geography, University of Zürich, Zürich, Switzerland

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970
Published: unpublished

License
Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

Abstract

This function module was written for Google Earth Engine (GEE) as an implementation of the Differential Evolution algorithm for optimizing functions (i.e., fitting curves) on remotely sensed imagery data. Its purpose is to allow the user to fit any arbitrary functional form on GEE's image collection objects, making the module particularly useful for time series analyses on satellite image collections that require flexible curve fitting algorithms (e.g., double-logistic functions with several parameters). The function makes extensive use of the array image object, which embeds multidimensional arrays at every pixel and thus allows for matrix calculations using pixel level time series matrices. Moreover, the module was designed to produce a variety of outputs: either a final value, i.e., an optimized set of parameters that can be used to fit a functional form or curve to the image collection data, or intermediate values called populations. Populations are lists of candidate parameters that are being considered in the algorithm for fitting a curve to the image collection data. The option to produce intermediate outputs in the form of full populations allows users to run analyses in series to refine the best fit possible, referred to as a "daisy chain" analysis when done in repetition. Moreover, producing intermediate outputs also allows users to run multiple populations in parallel. The function is implemented in both Javascript and Python, and an example on Sentinel-2 data time series is included.

Statement of Need

Optimization algorithms based on natural selection, also called genetic or evolutionary algorithms, have been used since the 1950's (Mitchell, 1998) and continue to be re-examined for use as well as development [Ahmad et al. (2022)](Das et al., 2016)[Das & Suganthan (2010)](Pant et al., 2020)(K. Price et al., 2006). The idea is simple: iterate a population of candidate solutions to a problem while mixing candidate solutions in the same ways that populations of organisms undergo genetic variation. For example, if you have observation points in 2 dimensions (see the algorithm figure below), and you need to fit a specific mathematical model (e.g., such as a logarithmic function including 3 parameters *a*, *b*, and *c*), the algorithm proceeds by (1) randomly generating a population of candidate mathematical models that are fit to the data (2) then undergoing an "evolution" process where you mix/combine best fitting models, iteratively, until a satisfactory model has been reached.

DRAFT

The Journal of Open Source Software

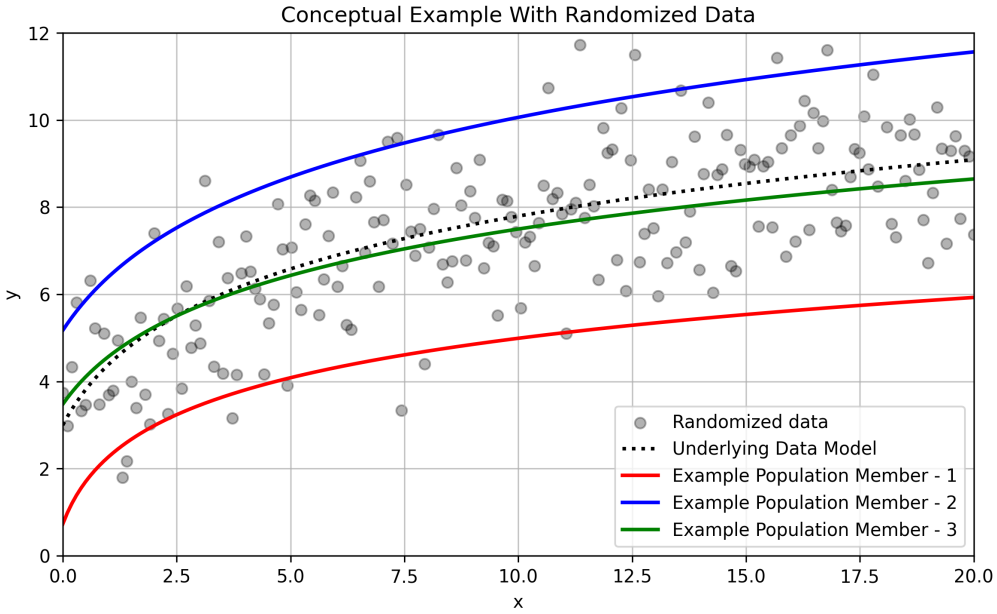

Figure 1: DE Optimization in a single chart

Storn and Price (Storn & Price, 1997) developed such an algorithm that they called differential evolution (DE) that is particularly suited for optimizing non-continuous and non-differentiable real-valued functions using real-valued parameters, and the approach has since been applied in a number of scholarly fields [Biesbroek (2006)](Feoktistov, 2006) [K. V. Price et al. (2005)](Qing, 2010). Nowadays, there are multiple software packages available to run DE, including a package in R (Mullen et al., 2011), a function in SciPy, and commercial tools in MatLab and Mathematica.

The implementation of DE in GEE aims to assist with modelling time series data using arbitrary mathematical models (e.g., linear, parabolic, logarithmic, etc.) on global-scale remote sensing datasets—especially when single pixels have sparse observations, such as when remotely sensed data is heavily impeded by cloud cover, and the process to fit a model through more standard regression modelling approaches are unsuitable. Moreover, the algorithm is structured to allow the production of intermediate outputs (in the form of populations of optimized models, rather than just a final optimized model), which allows users to take advantage of GEE's queueing system and gives greater options for task-based parallelization to accomplish computationally rigorous workflows on large amounts of imagery.

Implementation and Opportunities for Development

When searching for optimal parameter values, the algorithm benefits from a higher number of iterations in addition to a greater number of potential population members; i.e., optimization will improve when the algorithm is (1) testing more potential population candidates and (2) taking a greater number iteration steps to improve these potential options many times. Both a higher number of population members, as well as a greater number of iterations, require greater computational memory and resources. This implementation was structured accordingly to parallelize computation as much as possible while also making it possible to iteratively produce intermediary populations as evolution progresses.

More specifically, if users hit memory limits with their population number or their number of iterations, the algorithm allows users to structure their workflow via a divide-and-conquer approach: the users may run any number of populations independently of one another then combine the optimal parameter vectors from every separate population run into a final

DRAFT

population that can be further iterated. Dividing the populations in this way allows users to run multiple sub-populations in parallel as Earth Engine tasks, with each using the maximum amount of memory possible for a single task.

This means the maximum population size possible for this function is the maximum number of population candidates that can be run via a single Google Earth Engine task for 1 iteration; i.e., a job where all memory available is used for population size while still making a single iteration. In this case, a population of mutated vectors from a previous iteration is used as the input into a further iteration into a follow-up Google Earth Engine task. It's for this reason that the implementation allows the export of full populations, in the form of GEE array-images, rather than just single model parameter sets. It allows users to follow the daisy-chain procedure with the output of one iteration becoming the input into the following iteration.

The algorithm is furthermore structured to produce metrics that help the user decide whentheir time series model has been optimized to a desired degree of fitness; i.e., when a chosen RMSE value has been achieved. The implementation includes an option to produce an RMSE image, termed a *screelmage*, to monitor the progression of the optimization success with a scree plot similar to what is used in dimensional reduction techniques (Cattell, 1966). This image is comprised of multiple bands, wherein each band is the best RMSE value from the population at that iteration. It allows users to determine when convergence on an optimal value has been achieved and an acceptable final parameter set has been produced.

At the time of publication, the current mutation functions and the crossover functions are coded in an attempted modularized fashion so additional functions can be developed in the future. The current default mutation option *rand* randomly selects 3 parameter vectors from the population, computes the difference between 2 of the vectors, multiplies the difference by a scale factor (i.e., a real number between 0 and 2) then adds it to the 3rd randomly chosen vector; the other available mutation function (*best*) performs the same arithmetic except using the best parameter vector (according to RMSE) instead of a 3rd randomly chosen vector. The crossover function is a simple binomial wherein each iteration is tested according to whether a random number generated in the range of [0 to 1) from a uniform distribution is higher than a user supplied cross over value in the range of (0 to 1). Ongoing and future developments to the code include a helper function to randomly subsample dense input time series according to relative temporal density, allowing for greater control of the size of inputs so that memory limits can be better bypassed as well as potential crossover function variations.

Additional Functionality

In addition to the core functionality provided by the `de_optim` function, specific functions have been provided (both analytical and practical) to allow users the ability to customize their workflows.

Temporal Subsampling

Given the constraints on memory that effect the parameterization of the optimization process, a temporal subsampling function has been added to give users the ability to reduce the size of the image collections being used as the basis for their optimization tasks.

The subsampling function specifically allows for a temporal subsetting process, in which every pixel's time-series of observations is sampled according to their temporal density; i.e., time series are reduced to a specific size by removing the necessary number points at an individual pixel-level according to their relative frequency across the timespan through a weighted-random sampling process wherein the weights are the calculated as the density of observations around a specified time-window/kernel. See the documentation for more details.

DRAFT

Testing

Included in the repo is also a PyTest based testing framework to confirm the correct operation of the algorithm. It works by asserting specific algorithmic conditions when tested across arbitrary functional forms specified with randomly generated parameter sets. In other words, the testing process allows users to confirm that that algorithm:

- **Generates sets of functional coefficients** on **arbitrary** closed-form algebraic functions** such that each coefficient **falls within the numeric bounds** provided.
- Moreover, when the coefficient sets are being assessed **as the iterations proceed**, the fitness score of the coefficients **either improves or stabilizes without exception**.

The existing PyTest framework assesses a variety of functional families— currently including multiple replicates of logarithmic, exponential, harmonic, and linear functions—while also allowing users a structure to expand on the tests *ad hoc* in order to affirm algorithmic fidelity.

Acknowledgements

We would like to acknowledge that our research was funded partly by the Canton of Zürich and partly through a Google Earth Engine research award.

* Citations

Ahmad, M. F., Isa, N. A. M., Lim, W. H., & Ang, K. M. (2022). Differential evolution: A recent review based on state-of-the-art works. *Alexandria Engineering Journal*, *61*(5), 3831–3872. https://doi.org/10.1016/j.aej.2021.09.013

Biesbroek, R. (2006). A comparison of the differential evolution method with genetic algorithms for orbit optimisation. *57th International Astronautical Congress*, C1–4. https://doi.org/10.2514/6.iac-06-c1.4.02

Cattell, R. B. (1966). The scree test for the number of factors. *Multivariate Behavioral Research*, *1*(2), 245–276. https://doi.org/10.1207/s15327906mbr0102_10

Das, S., Mullick, S. S., & Suganthan, P. N. (2016). Recent advances in differential evolution–an updated survey. *Swarm and Evolutionary Computation*, *27*, 1–30. https://doi.org/10.1016/j.swevo.2016.01.004

Das, S., & Suganthan, P. N. (2010). Differential evolution: A survey of the state-of-the-art. *IEEE Transactions on Evolutionary Computation*, *15*(1), 4–31. https://doi.org/10.1016/j.aej.2021.09.013

Feoktistov, V. (2006). *Differential evolution*. Springer. https://doi.org/10.1007/978-0-387-36896-2

Mitchell, M. (1998). *An introduction to genetic algorithms*. MIT press. https://doi.org/10.7551/mitpress/3927.001.0001

Mullen, K. M., Ardia, D., Gil, D. L., Windover, D., & Cline, J. (2011). DEoptim: An r package for global optimization by differential evolution. *Journal of Statistical Software*, *40*, 1–26. https://doi.org/10.32614/CRAN.package.DEoptim

Pant, M., Zaheer, H., Garcia-Hernandez, L., Abraham, A., & others. (2020). Differential evolution: A review of more than two decades of research. *Engineering Applications of Artificial Intelligence*, *90*, 103479. https://doi.org/10.1016/j.engappai.2020.103479

Price, K. V., Storn, R. M., Lampinen, J. A., Wormington, M., Matney, K. M., & Bowen, D. K. (2005). Application of differential evolution to the analysis of x-ray reflectivity data. *Differential Evolution: A Practical Approach to Global Optimization*, 463–478. https://doi.org/10.1007/3-540-31306-0_16

DRAFT

Price, K., Storn, R. M., & Lampinen, J. A. (2006). *Differential evolution: A practical approach to global optimization*. Springer Science & Business Media. https://doi.org/10.1007/3-540-31306-0

Qing, A. (2010). Basics of differential evolution. In *Differential evolution in electromagnetics* (pp. 19–42). Springer. https://doi.org/10.1007/978-3-642-12869-1_2

Storn, R., & Price, K. (1997). Differential evolution–a simple and efficient heuristic for global optimization over continuous spaces. *Journal of Global Optimization*, *11*, 341–359. https://doi.org/10.1023/A:1008202821328

DRAFT